

demo

January 13, 2024

```
[1]: import numpy as np
import pandas as pd
import pickle

from modules.logistic_regression_cuda import LogisticRegression
from modules.nn import NNEstimator
from modules.svm import MultiClassSVM
from modules.pca import PCA
from modules.model_selection import CrossValidate, GridSearch, BootStrap

[2]: # load data
with open('../data/train_feature.pkl', 'rb') as f:
    dataset_feature = pickle.load(f)
    dataset_feature = np.array(dataset_feature.todense()) # (11314, 10000)
dataset_label = np.load('../data/train_labels.npy') # (11314,)

with open('../data/test_feature.pkl', 'rb') as f:
    test_feature = pickle.load(f)
    test_feature = np.array(test_feature.todense()) # (7532, 10000)
```

0.1 SVM

```
[3]: lambda_ = 1 / 1e7
svm = MultiClassSVM(C=1/lambda_, lr=lambda_)
svm.fit(dataset_feature, dataset_label, n_jobs=20)

[4]: svm_pred = svm.predict(test_feature)
pd.DataFrame({'ID': np.arange(0, len(svm_pred)), 'label': svm_pred}).
    to_csv('submit_SVM.csv', index=False)
```

0.2 Logistic Regression

We use CUDA to accelerate the training process by default. If CUDA is not available, use `from

```
[5]: # if cuda is not available:
# from modules.logistic_regression import LogisticRegression
```

```
lr = LogisticRegression(lr=100, n_iters=1000)
lr.fit(dataset_feature, dataset_label)
print("train accuracy: ", lr.score(dataset_feature, dataset_label))
```

train accuracy: 0.9990277532260916

```
[6]: lr_pred = lr.predict(test_feature)
pd.DataFrame({'ID': np.arange(0, len(lr_pred)), 'label': lr_pred}).
    to_csv('submit_LR.csv', index=False)
```

0.3 Nerual Network

```
[7]: nn_net = NNEstimator(weight_decay=1e-7,hidden_size=1024,drop_rate=0.
    to 9,lr=1e-3,epoch_num=20)
nn_net.fit(dataset_feature, dataset_label)
print("train accuracy: ", nn_net.score(dataset_feature, dataset_label))
```

train accuracy: 0.9979671203818278

```
[8]: nn_pred = nn_net.predict(test_feature)
pd.DataFrame({'ID': np.arange(0, len(nn_pred)), 'label': nn_pred}).
    to_csv('submit_NN.csv', index=False)
```

0.4 Model Selection

```
[9]: ## Cross Validate

cv = CrossValidate(LogisticRegression(lr=100, n_iters=1000),n_folds=5)
cv.fit(dataset_feature, dataset_label)
print(cv.get_result())
```

[0.8926204153778171, 0.9058771542200619, 0.9054352629253204, 0.900574458683164, 0.9022988505747126]

```
[10]: ## BootStrap

bs = BootStrap(LogisticRegression(lr=100, n_iters=1000),n_folds=5)
bs.fit(dataset_feature, dataset_label)
print(bs.get_result())
```

[0.8847590953785645, 0.8876315267024022, 0.8851113716295428, 0.8818882085719929, 0.8789958815454011]

0.5 Model Search

```
[11]: param_grid = {
        'lr': [1e-4, 1e-5, 1e-6, 1e-7],
        'C': [1e4, 1e5, 1e6, 1e7],
    }
    gs_svm = GridSearch(MultiClassSVM(), param_grid, 5)
    gs_svm.fit(dataset_feature, dataset_label, n_jobs=20)
    gs_svm_results_df = pd.DataFrame(columns=list(param_grid.keys())+['score'])
    for params, score in gs_svm.results:
        gs_svm_results_df = gs_svm_results_df.append(pd.Series(**params, 'score':
↪score)), ignore_index=True)
    gs_svm_results_df
```

```
[11]:
```

	lr	C	score
0	1.000000e-04	10000.0	0.773285
1	1.000000e-04	100000.0	0.726169
2	1.000000e-04	1000000.0	0.585468
3	1.000000e-04	10000000.0	0.585295
4	1.000000e-05	10000.0	0.883950
5	1.000000e-05	100000.0	0.865739
6	1.000000e-05	1000000.0	0.851602
7	1.000000e-05	10000000.0	0.844354
8	1.000000e-06	10000.0	0.895173
9	1.000000e-06	100000.0	0.891109
10	1.000000e-06	1000000.0	0.888898
11	1.000000e-06	10000000.0	0.868130
12	1.000000e-07	10000.0	0.861942
13	1.000000e-07	100000.0	0.903836
14	1.000000e-07	1000000.0	0.899417
15	1.000000e-07	10000000.0	0.885539

```
[12]: param_grid = {
        'lr': [1e-3],
        'drop_rate': [0.7, 0.8, 0.9],
        'hidden_size': [1024, 1024+512, 2048,],
        'epoch_num': [15, 20, 25],
        'weight_decay': [0, 1e-7, 1e-6, 1e-5]
    }
    gs_nn = GridSearch(NNEstimator(), param_grid, 5)
    gs_nn.fit(dataset_feature, dataset_label)
    gs_nn_results_df = pd.DataFrame(columns=list(param_grid.keys())+['score'])
    for params, score in gs_nn.results:
        gs_nn_results_df = gs_nn_results_df.append(pd.Series(**params, 'score':
↪score)), ignore_index=True)
    gs_nn_results_df
```

```
[12]:
```

	lr	drop_rate	hidden_size	epoch_num	weight_decay	score
0	0.001	0.7	1024.0	15.0	0.000000e+00	0.911614
1	0.001	0.7	1024.0	15.0	1.000000e-07	0.912056
2	0.001	0.7	1024.0	15.0	1.000000e-06	0.912675
3	0.001	0.7	1024.0	15.0	1.000000e-05	0.910377
4	0.001	0.7	1024.0	20.0	0.000000e+00	0.913470
..
103	0.001	0.9	2048.0	20.0	1.000000e-05	0.911437
104	0.001	0.9	2048.0	25.0	0.000000e+00	0.911172
105	0.001	0.9	2048.0	25.0	1.000000e-07	0.915768
106	0.001	0.9	2048.0	25.0	1.000000e-06	0.913381
107	0.001	0.9	2048.0	25.0	1.000000e-05	0.908786

[108 rows x 6 columns]

```
[13]: param_grid = {
        'lr': [0.1, 1, 10, 100, 1000],
        'n_iters': [1000, 2000],
    }
gs_lr = GridSearch(LogisticRegression(), param_grid, 5)
gs_lr.fit(dataset_feature, dataset_label)
gs_lr_results_df = pd.DataFrame(columns=list(param_grid.keys())+['score'])
for params, score in gs_lr.results:
    gs_lr_results_df = gs_lr_results_df.append(pd.Series(**params, 'score':
↪score)), ignore_index=True)
gs_lr_results_df
```

```
[13]:
```

	lr	n_iters	score
0	0.1	1000.0	0.669521
1	0.1	2000.0	0.733162
2	1.0	1000.0	0.822520
3	1.0	2000.0	0.841701
4	10.0	1000.0	0.885539
5	10.0	2000.0	0.892522
6	100.0	1000.0	0.898886
7	100.0	2000.0	0.902599
8	1000.0	1000.0	0.880414
9	1000.0	2000.0	0.884480

0.6 PCA

```
[14]: pca = PCA(n_components=5000)
pca.fit(dataset_feature)
```

```
[15]: dataset_feature_pca = pca.transform(dataset_feature)
test_feature_pca = pca.transform(test_feature)
```

```
[16]: lr_pca = LogisticRegression(lr=100, n_iters=1000)
lr_pca.fit(dataset_feature_pca, dataset_label)
print("train accuracy: ", lr_pca.score(dataset_feature_pca, dataset_label))
lr_pred_pca = lr_pca.predict(test_feature_pca)
pd.DataFrame({'ID': np.arange(0, len(lr_pred_pca)), 'label': lr_pred_pca}).
  to_csv('submit_LR_pca.csv', index=False)
```

train accuracy: 0.9949619939897472

```
[17]: svm_pca = MultiClassSVM(C=1e5, lr=1e-7)
svm_pca.fit(dataset_feature_pca, dataset_label, n_jobs=20)
print("train accuracy: ", svm_pca.score(dataset_feature_pca, dataset_label))
svm_pred_pca = svm_pca.predict(test_feature_pca)
pd.DataFrame({'ID': np.arange(0, len(svm_pred_pca)), 'label': svm_pred_pca}).
  to_csv('submit_SVM_pca.csv', index=False)
```

train accuracy: 0.9642920275764539

```
[18]: nn_pca = NNEstimator(weight_decay=1e-7, hidden_size=1024, drop_rate=0.
  9, lr=1e-3, epoch_num=20)
nn_pca.fit(dataset_feature_pca, dataset_label)
print("train accuracy: ", nn_pca.score(dataset_feature_pca, dataset_label))
nn_pred_pca = nn_pca.predict(test_feature_pca)
pd.DataFrame({'ID': np.arange(0, len(nn_pred_pca)), 'label': nn_pred_pca}).
  to_csv('submit_NN_pca.csv', index=False)
```

train accuracy: 0.9950503800601025